

# Exploratory Data Analysis

## Telco Customer Churn

### ABOUT DATASET

**gender** - Whether the customer is a male or a female

**SeniorCitizen** - Whether the customer is a senior citizen or not (1, 0)

**Partner** - Whether the customer has a partner or not (Yes, No)

**Dependents** - Whether the customer has dependents or not (Yes, No)

**Tenure** - Number of months the customer has stayed with the company check

**PhoneService** - Whether the customer has a phone service or not (Yes, No)

**MultipleLines** - Whether the customer has multiple lines or not (Yes, No, No phone service)

**InternetService** - Customer's internet service provider (DSL, Fiber optic, No)

**OnlineSecurity** - Whether the customer has online security or not (Yes, No, No internet service)

**OnlineBackup** - Whether the customer has online backup or not (Yes, No, No internet service)

**DeviceProtection** - Whether the customer has device protection or not (Yes, No, No internet service)

**TechSupport** - Whether the customer has tech support or not (Yes, No, No internet service)

**StreamingTV** - Whether the customer has streaming TV or not (Yes, No, No internet service)

**StreamingMovies** - Whether the customer has streaming movies or not (Yes, No, No internet service)

**Contract** - The contract term of the customer (Month-to-month, One year, Two year)

**PaperlessBilling** - Whether the customer has paperless billing or not (Yes, No)

**PaymentMethod** - The customer's payment method (Electronic check, Mailed check, Bank transfer (automatic), Credit card (automatic))

**MonthlyCharges** - The amount charged to the customer monthly

**TotalCharges** - The total amount charged to the customer

**Churn** - Whether the customer churned or not (Yes or No)

Churn is the number of customers who left or the number of customers who cancel or don't renew a subscription. Churn is target variable or dependent variables in our dataset.

### 1) Understanding the variables

First, we need to understanding all the variables in our dataset. It is crucial because it can provides important insights into the data and helps us make informed decisions during the next analysis process.

```
library(readr)
df = read_csv("/Users/shafaqonitatingmail.com/Documents/Semester 4/Data Mining and
Visualization/Assignment/LAB/Telco Customer Churn/WA_Fn-UseC_-Telco-Customer-
Churn.csv", show_col_types = FALSE)
```

By adding `show_col_types = FALSE`, we can eliminate the warning message that appears if the detected column type does not match the given definition.

```
head(df)
```

A tibble: 6 × 21

customerID <chr>	gender <chr>	SeniorCitizen <dbl>	Partner <chr>	Dependents <chr>	tenure <dbl>	PhoneService <chr>	MultipleLines <chr>	InternetService <chr>	
7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	
5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	
3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	
7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	
9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	
9305-CDSKC	Female	0	No	No	8	Yes	Yes	Fiber optic	

6 rows | 1-9 of 21 columns

From the `head()` function above, we can see that of the 21 columns we have, almost all of them are of character type and three of them are of double type.

```
str(df)
```

```
spc_tbl_ [7,043 × 21] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ customerID      : chr [1:7043] "7590-VHVEG" "5575-GNVDE" "3668-QPYBK" "7795-CFOCW" ...
 $ gender          : chr [1:7043] "Female" "Male" "Male" "Male" ...
 $ SeniorCitizen   : num [1:7043] 0 0 0 0 0 0 0 0 0 ...
 $ Partner         : chr [1:7043] "Yes" "No" "No" "No" ...
 $ Dependents      : chr [1:7043] "No" "No" "No" "No" ...
 $ tenure          : num [1:7043] 1 34 2 45 2 8 22 10 28 62 ...
 $ PhoneService    : chr [1:7043] "No" "Yes" "Yes" "No" ...
 $ MultipleLines   : chr [1:7043] "No phone service" "No" "No" "No phone service" ...
 $ InternetService : chr [1:7043] "DSL" "DSL" "DSL" "DSL" ...
 $ OnlineSecurity  : chr [1:7043] "No" "Yes" "Yes" "Yes" ...
 $ OnlineBackup    : chr [1:7043] "Yes" "No" "Yes" "No" ...
 $ DeviceProtection: chr [1:7043] "No" "Yes" "No" "Yes" ...
 $ TechSupport     : chr [1:7043] "No" "No" "No" "Yes" ...
 $ StreamingTV     : chr [1:7043] "No" "No" "No" "No" ...
 $ StreamingMovies : chr [1:7043] "No" "No" "No" "No" ...
 $ Contract        : chr [1:7043] "Month-to-month" "One year" "Month-to-month" "One year" ...
 $ PaperlessBilling: chr [1:7043] "Yes" "No" "Yes" "No" ...
 $ PaymentMethod   : chr [1:7043] "Electronic check" "Mailed check" "Mailed check" "Bank transfer (automatic)" ...
 $ MonthlyCharges  : num [1:7043] 29.9 57 53.9 42.3 70.7 ...
 $ TotalCharges    : num [1:7043] 29.9 1889.5 108.2 1840.8 151.7 ...
 $ Churn           : chr [1:7043] "No" "No" "Yes" "No" ...
```

From `str()` function, we can see the dimension of our data (*row vs columns*) and data types for each variables. From the output above we can see that we have 7043 observations and 21 variables and data types for each variables, which character for `customerID`, `gender`, `Partner`, `Dependents`, `PhoneService`, `MultipleLines`, `InternetService`, `OnlineSecurity`, `OnlineBackup`, `DeviceProtection`, `TechSupport`, `StreamingTV`, `StreamingMovies`, `Contract`, `PaperlessBilling`, `PaymentMethod`, `Churn` and double for `SeniorCitizen`, `tenure`, `MonthlyCharges` and `TotalCharges`.

```
summary(df)
```

customerID Length:7043 Class :character Mode :character	gender Length:7043 Class :character Mode :character	SeniorCitizen Min. : 0.0000 1st Qu.:0.0000 Median :0.0000 Mean :0.1621 3rd Qu.:0.0000 Max. :1.0000	Partner Length:7043 Class :character Mode :character	Dependents Length:7043 Class :character Mode :character	tenure Min. : 0.00 1st Qu.: 9.00 Median :29.00 Mean :32.37 3rd Qu.:55.00 Max. :72.00	PhoneService Length:7043 Class :character Mode :character	MultipleLines Length:7043 Class :character Mode :character
InternetService Length:7043 Class :character Mode :character	OnlineSecurity Length:7043 Class :character Mode :character	OnlineBackup Length:7043 Class :character Mode :character	DeviceProtection Length:7043 Class :character Mode :character	TechSupport Length:7043 Class :character Mode :character	StreamingTV Length:7043 Class :character Mode :character	StreamingMovies Length:7043 Class :character Mode :character	Contract Length:7043 Class :character Mode :character
PaperlessBilling Length:7043 Class :character Mode :character	PaymentMethod Length:7043 Class :character Mode :character	MonthlyCharges Min. : 18.25 1st Qu.: 35.50 Median : 70.35 Mean : 64.76 3rd Qu.: 89.85 Max. :118.75	TotalCharges Min. : 18.8 1st Qu.: 401.4 Median :1397.5 Mean :2283.3 3rd Qu.:3794.7 Max. :8684.8 NA's :11	Churn Length:7043 Class :character Mode :character			

From the output above, we can see that there are 11 null values in the `TotalCharges` column and there are many character data types that we cannot see the minimum, quarter, and maximum values so we can only conclude that there are no null values in it.

We know that almost all of the character data types in the dataset are categorical data types with the output that it produces. Now, we check whether there are misformatted data using `unique()` function to categorical variables we have

```
print(unique(df$gender))
print(unique(df$Partner))
print(unique(df$Dependents))
print(unique(df$PhoneService))
print(unique(df$MultipleLines))
print(unique(df$InternetService))
print(unique(df$OnlineSecurity))
print(unique(df$OnlineBackup))
print(unique(df$DeviceProtection))
print(unique(df$TechSupport))
print(unique(df$StreamingTV))
print(unique(df$StreamingMovies))
print(unique(df$Contract))
print(unique(df$PaperlessBilling))
print(unique(df$PaymentMethod))
print(unique(df$Churn))
```

```
[1] "Female" "Male"
[1] "Yes" "No"
[1] "No" "Yes"
[1] "No" "Yes"
[1] "No phone service" "No" "Yes"
[1] "DSL" "Fiber optic" "No"
[1] "No" "Yes" "No internet service"
[1] "Yes" "No" "No internet service"
[1] "No" "Yes" "No internet service"
[1] "No" "Yes" "No internet service"
[1] "No" "Yes" "No internet service"
[1] "No" "Yes" "No internet service"
[1] "Month-to-month" "One year" "Two year"
[1] "Yes" "No"
[1] "Electronic check" "Mailed check" "Bank transfer (automatic)" "Credit card (automatic)"
[1] "No" "Yes"
```

Based on the output, there are no categorical variables that have misformatted data or in other words the output of each categorical data match the data that should be on about dataset.

## 2) Cleaning dataset

Now, we delete our `CustomerID` column because we don't need it to do data analysis in the future or in other words, there is no significant information that we will get if we keep the column.

```
df = subset(df, select = -c(customerID))
```

```
head(df)
```

A tibble: 6 × 20

gender <chr>	SeniorCitizen <dbl>	Partner <chr>	Dependents <chr>	tenure <dbl>	PhoneService <chr>	MultipleLines <chr>	InternetService <chr>	OnlineSecurity <chr>	OnlineBackup <chr>	
Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	
Male	0	No	No	34	Yes	No	DSL	Yes	No	
Male	0	No	No	2	Yes	No	DSL	Yes	Yes	
Male	0	No	No	45	No	No phone service	DSL	Yes	No	
Female	0	No	No	2	Yes	No	Fiber optic	No	No	
Female	0	No	No	8	Yes	Yes	Fiber optic	No	No	

6 rows | 1–10 of 20 columns

From the output above, we can see that there is no longer a CustomerID column in our dataframe.

To synchronize all the categorical data types we have, we need to change the output produced by SeniorCitizen because if we notice the output is different from our other categorical outputs, which are numbers (for example: 0, 1, or 2 (for more than 2 outputs))

```
df$SeniorCitizen <- ifelse(df$SeniorCitizen == 1, "Yes", "No")
```

head(df)

A tibble: 6 × 20

gender <chr>	SeniorCitizen <chr>	Partner <chr>	Dependents <chr>	tenure <dbl>	PhoneService <chr>	MultipleLines <chr>	InternetService <chr>	OnlineSecurity <chr>	OnlineBackup <chr>	
Female	No	Yes	No	1	No	No phone service	DSL	No	Yes	
Male	No	No	No	34	Yes	No	DSL	Yes	No	
Male	No	No	No	2	Yes	No	DSL	Yes	Yes	
Male	No	No	No	45	No	No phone service	DSL	Yes	No	
Female	No	No	No	2	Yes	No	Fiber optic	No	No	
Female	No	No	No	8	Yes	Yes	Fiber optic	No	No	

6 rows | 1–10 of 20 columns

From the output above, we can see that on the SeniorCitizen column, the output has changed to a numeric output, mainly binary (0 or 1)

In accordance with part **1) Understanding the variables**, there are 11 NAs in totalCharges. Now, we can solve this problem, I choose to eliminate them because those 11 null values are only a small part of the whole data set (less than 1% of the total data), eliminating them is practically insignificant in terms of information loss or bias.

```
df <- na.omit(df)
```

summary(df)

gender Length:7032 Class :character Mode :character	SeniorCitizen Length:7032 Class :character Mode :character	Partner Length:7032 Class :character Mode :character	Dependents Length:7032 Class :character Mode :character	tenure Min. : 1.00 1st Qu.: 9.00 Median :29.00 Mean :32.42 3rd Qu.:55.00 Max. :72.00	PhoneService Length:7032 Class :character Mode :character	MultipleLines Length:7032 Class :character Mode :character	InternetService Length:7032 Class :character Mode :character
OnlineSecurity Length:7032 Class :character Mode :character	OnlineBackup Length:7032 Class :character Mode :character	DeviceProtection Length:7032 Class :character Mode :character	TechSupport Length:7032 Class :character Mode :character	StreamingTV Length:7032 Class :character Mode :character	StreamingMovies Length:7032 Class :character Mode :character	Contract Length:7032 Class :character Mode :character	PaperlessBilling Length:7032 Class :character Mode :character
PaymentMethod Length:7032 Class :character Mode :character	MonthlyCharges Min. : 18.25 1st Qu.: 35.59 Median : 70.35 Mean : 64.80 3rd Qu.: 89.86 Max. :118.75	TotalCharges Min. : 18.8 1st Qu.: 401.4 Median :1397.5 Mean :2283.3 3rd Qu.:3794.7 Max. :8684.8	Churn Length:7032 Class :character Mode :character				

From the output above, we can now see that there are no more null values in our dataset.


Now using boxplots, we check if there are any outliers in our dataset, especially in the numerical variables, namely Tenure, MonthlyCharges, and TotalCharges.

```
# outliers
b1 = ggplot(data = df, aes(x = "", y = tenure)) +
  geom_boxplot(fill = "lightblue", color = "#142850") +
  labs(x = "", y = "Tenure", title = "Boxplot of Tenure")


b2 = ggplot(data = df, aes(x = "", y = MonthlyCharges)) +
  geom_boxplot(fill = "lightblue", color = "#142850") +
  labs(x = "", y = "MonthlyCharges", title = "Boxplot of Monthly Charges")

b3 = ggplot(data = df, aes(x = "", y = TotalCharges)) +
  geom_boxplot(fill = "lightblue", color = "#142850") +
  labs(x = "", y = "TotalCharges", title = "Boxplot of Total Charges")


b1
b2
b3
```



From the boxplot on the left, we can conclude that **there is no outliers in Tenure**. It means that there are no data points that are very far from the rest of the data that could indicate potential anomalies or extreme values.



And also, from the boxplot on the left, we can conclude that **there is no outliers in Monthly Charges**. It means that there are no data points that are very far from the rest of the data that could indicate potential anomalies or extreme values.



And the last is from TotalCharges, we can conclude that **there is no outliers in Total Charges**. It means that there are no data points that are very far from the rest of the data that could indicate potential anomalies or extreme values.

### 3) Visualizing and Analyzing Results

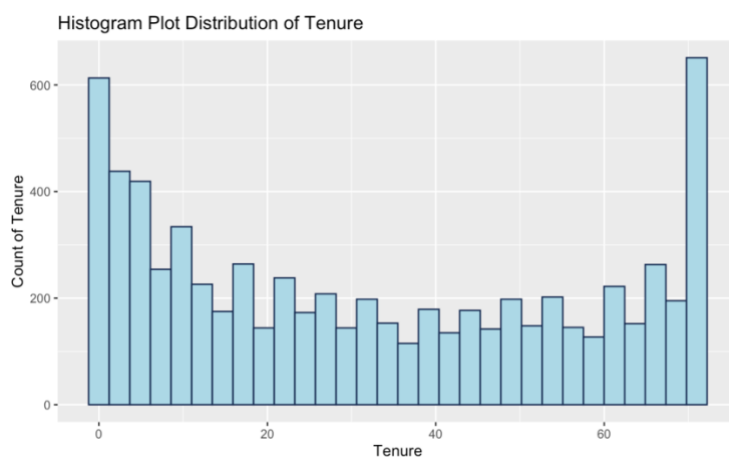
Now, we perform a histogram to show the distribution plot of the numerical variable. In our dataset are Tenure, MonthlyCharges, and TotalCharges

```
# Histogram Plot Distribution of tenure
h1 = ggplot(data = df, aes(x = tenure)) +
  geom_histogram(fill = "lightblue", color = "#142850", bins = 30) +
  labs(x = "Tenure", y = "Count of Tenure", title = "Histogram Plot Distribution of Tenure")


# Histogram Plot Distribution of MonthlyCharges
h2 = ggplot(data = df, aes(x = MonthlyCharges)) +
  geom_histogram(fill = "lightblue", color = "#142850", bins = 30) +
  labs(x = "MonthlyCharges", y = "Count of MonthlyCharges", title = "Histogram Plot Distribution of MonthlyCharges")

# Histogram Plot Distribution of TotalCharges
h3 = ggplot(data = df, aes(x = TotalCharges)) +
  geom_histogram(fill = "lightblue", color = "#142850", bins = 30) +
  labs(x = "TotalCharges", y = "Count of TotalCharges", title = "Histogram Plot Distribution of TotalCharges")

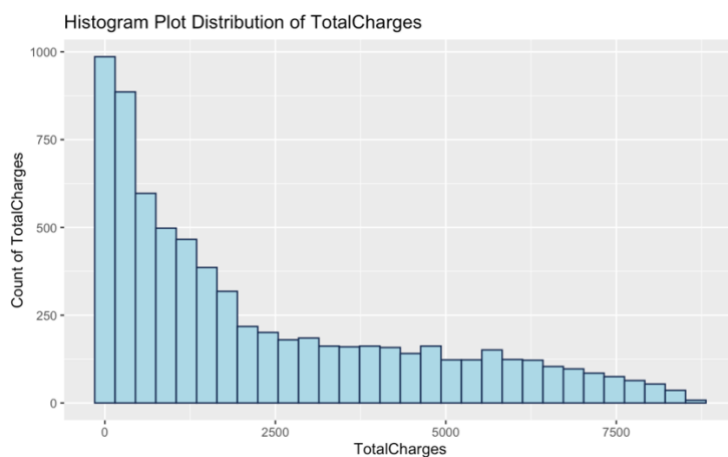
h1
h2
h3
```



From the histogram plot on the left, we can conclude that the number of months customers stay with the company is the most, which is more than 60 months and the second is 0 months. So, among the many customers, **the company has the most customers who have stayed more than 60 months** compared to other customers.



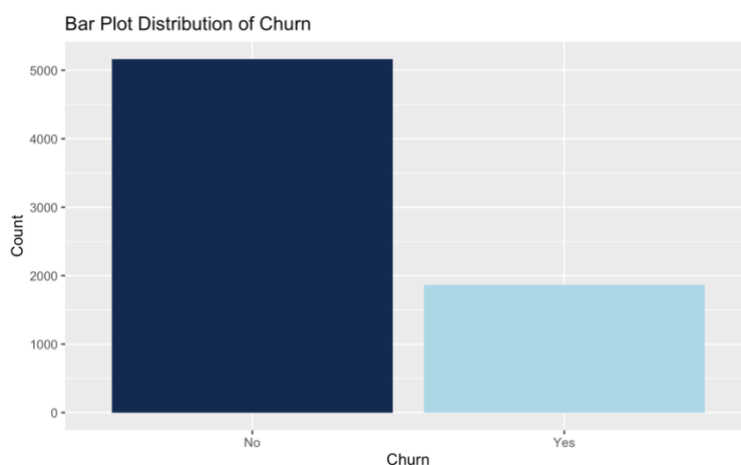
From the histogram plot on the left, we can conclude that the amount charged to the customer monthly is the most less than 25\$ (there is no information about the unit so I'll go with dollars). So, among the many customers, **the company has the most customers who get a monthly charges of less than \$25.**



From the histogram plot on the left, we can infer that the total amount charged to customers during the subscription is mostly less than \$0 (there is no information about the unit, so I will use dollars). So, among the many customers, **the company has the most customers who get total charges less than \$0 and the more expensive the total charges are, the less customers pay the total charges.**

Now, we use bar plot to see the distribution of our Churn data

```
ggplot(data = df, aes(x = Churn)) +
  geom_bar(fill = c("#142850", "lightblue")) +
  labs(x = "Churn", y = "Count", title = "Bar Plot Distribution of Churn")
```



From the bar plot on the left, we can conclude that **the company has the most customers who are not churned or stayed with them or do not leave the company.**

Now we want to know:

1. What type of account services do customers (who are churned and not) have?
2. What type of internet services do customers (who are churned and not) have?
3. What type of citizen do the company have the most?
4. What gender of customers do the company have the most?

To answers those questions we can use bar plot with Churn grouped

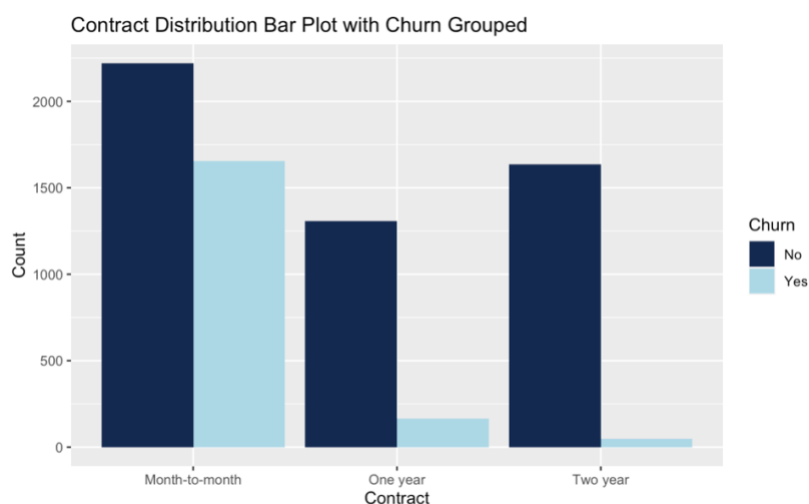
```
# Contract Distribution Bar Plot with Churn Grouped
g1 <- ggplot(data = df, aes(x = Contract)) +
  geom_bar(aes(fill = Churn), position = "dodge") +
  labs(x = "Contract", y = "Count", title = "Contract Distribution Bar Plot with
Churn Grouped") +
  scale_fill_manual(values = c("#142850", "lightblue"))

# Internet Service Bar Plot with Churn Grouped
g2 <- ggplot(data = df, aes(x = InternetService)) +
  geom_bar(aes(fill = Churn), position = "dodge") +
  labs(x = "Internet Service", y = "Count", title = " Internet Service Bar Plot
with Churn Grouped") +
  scale_fill_manual(values = c("#142850", "lightblue"))

# Internet Service Bar Plot with Churn Grouped
g3 <- ggplot(data = df, aes(x = SeniorCitizen)) +
  geom_bar(aes(fill = Churn), position = "dodge") +
  labs(x = "Senior Citizen", y = "Count", title = "Senior Citizen Distribution Bar
Plot with Churn Grouped") +
  scale_fill_manual(values = c("#142850", "lightblue"))

g1
g2
g3
```

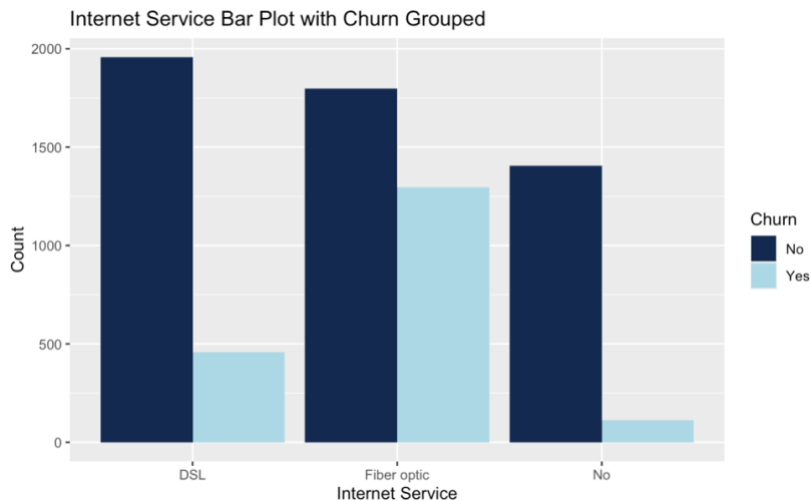
## 1. What type of i services do customers (who are churned and not) have?



From the plot on the left, we know that there are three contract, which Month-to-Month, One year, and Two year. We can conclude that **Month-to-month contract or service is the most favorite among customers who stay and are also the contract or service where many customers leave the company.**


## 2. What type of internet services do customers (who are churned and not) have?






From the plot on the left, we know that there are three internet services that customer use, which DSL, Fiber optic, and No (It means that the customers don't use internet services). We can conclude that **DSL is the most favorite internet services among customers who stay.** And **Fiber optic is the internet services where many customers leave the company.**

### 3. What type of customers do the company have the most?



From the plot on the left, we know that there are two citizen types, which senior citizen and not senior citizen. We can conclude that **most of the customers who stay with the company are not the senior citizens.**

### 4. What gender of customers do the company have the most?



From the plot on the left, we can conclude that **the company has more male customers who stay with the company compared to female customers**, where the difference is not that significant. And there is also not significant difference between the male and female customers who have leave the company.

# Predictive Modelling

## Telco Customer Churn

### 1) Data Pre-Processing

Before we do the predictive modelling, we have to do the Data Pre-Processing, we use oversampling method to address class imbalance in a Churn column. Class imbalance can lead to biased models that perform poorly in predicting the minority class.

```
library(caret)

# convert Churn variable to factor
df$Churn <- as.factor(df$Churn)

# perform oversampling
dfb <- upSample(x = df[, -which(names(df) == "Churn")],
               y = df$Churn,
               yname = "Churn")

# check the balanced class distribution
table(dfb$Churn)
```

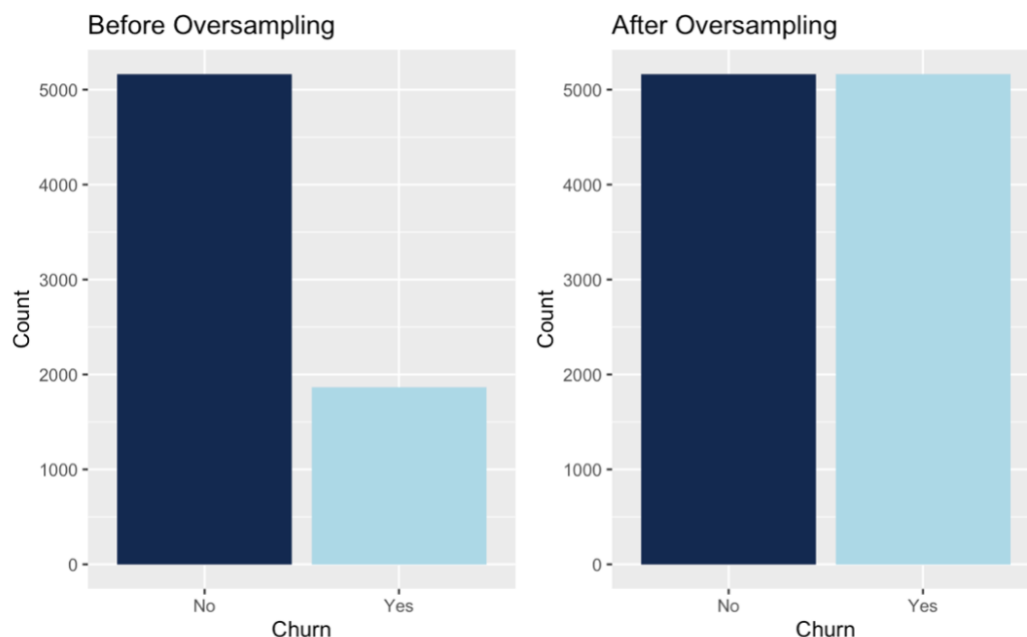
No	Yes
5163	5163

Now we have a balanced Churn column and we put it inside our new data frame, called dfb (short for data frame balance). And now, we plot the difference between the Churn column before oversampling and after oversampling.

```
# Before oversampling
b1 = ggplot(data = df, aes(x = Churn)) +
  geom_bar(fill = c("#142850", "lightblue")) +
  labs(x = "Churn", y = "Count", title = "Before Oversampling")

# Create a bar plot for the distribution of Churn in balanced_df
b2 = ggplot(data = dfb, aes(x = Churn)) +
  geom_bar(fill = c("#142850", "lightblue")) +
  labs(x = "Churn", y = "Count", title = "After Oversampling")

library(gridExtra)
grid.arrange(b1, b2, ncol=2)
```



The next step is to transform the format of our categorical variables into factors using the `factor()` function as it is required to perform a classification model using Random Forest. Because in classification modelling, it is important to state that the variable is a factor, so the algorithm will not treat the variable as an ordinal or interval and will treat the variable as a categorical variable and prevent unwanted bias.

```
dfb$gender <- factor(dfb$gender)
dfb$SeniorCitizen <- factor(dfb$SeniorCitizen)
dfb$Partner <- factor(dfb$Partner)
dfb$Dependents <- factor(dfb$Dependents)
dfb$PhoneService <- factor(dfb$PhoneService)
dfb$MultipleLines <- factor(dfb$MultipleLines)
dfb$InternetService <- factor(dfb$InternetService)
dfb$OnlineSecurity <- factor(dfb$OnlineSecurity)
dfb$OnlineBackup <- factor(dfb$OnlineBackup)
dfb$DeviceProtection <- factor(dfb$DeviceProtection)
dfb$TechSupport <- factor(dfb$TechSupport)
dfb$StreamingTV <- factor(dfb$StreamingTV)
dfb$StreamingMovies <- factor(dfb$StreamingMovies)
dfb$Contract <- factor(dfb$Contract)
dfb$PaperlessBilling <- factor(dfb$PaperlessBilling)
dfb$PaymentMethod <- factor(dfb$PaymentMethod)
dfb$Churn <- factor(dfb$Churn)
```

```
head(dfb)
```

Description: df [6 x 20]

	gender <fctr>	SeniorCitizen <fctr>	Partner <fctr>	Dependents <fctr>	tenure <dbl>	PhoneService <fctr>	MultipleLines <fctr>	InternetService <fctr>	OnlineSecurity <fctr>
1	Female	No	Yes	No	1	No	No phone service	DSL	No
2	Male	No	No	No	34	Yes	No	DSL	Yes
3	Male	No	No	No	45	No	No phone service	DSL	Yes
4	Male	No	No	Yes	22	Yes	Yes	Fiber optic	No
5	Female	No	No	No	10	No	No phone service	DSL	Yes
6	Male	No	No	Yes	62	Yes	No	DSL	Yes

6 rows | 1-10 of 20 columns

From the output above, our categorical data was successfully converted into factor types indicated by `<fctr>` under the column headings.

## 2) Modelling

**First step:** Split dataset into training and testing sets and create

```
# split the data into training and testing sets
library(caTools)
set.seed(1)
split <- sample.split(dfb$Churn, SplitRatio = 0.8)
train <- subset(dfb, split == TRUE)
test <- subset(dfb, split == FALSE)
```

We use a split ratio = 0.8, meaning we split our data frame 80% for the training set and 20% for the test set.

Then we create a formula of our dependent and independent variables that will be used in our model, we excluded variable Churn because it used as our dependent variable

```
n <- names(train)
excluded_vars <- c("Churn")
f <- as.formula(paste("Churn ~", paste(n[!n %in% excluded_vars], collapse = " +
")))
```

**Second step:** Create model and Train the model using our training set

```
library(randomForest)
modelRF <- randomForest(f, data=train)
print(modelRF)
```

Call:

```
randomForest(formula = f, data = train)
      Type of random forest: classification
      Number of trees: 500
```

No. of variables tried at each split: 4

OOB estimate of error rate: 11.65%

Confusion matrix:

	No	Yes	class.error
No	3387	743	0.17990315
Yes	219	3911	0.05302663

From the output above, we know that our modelling is use a classification. The number of trees is 500, meaning our model use 500 decision trees and it is as the default value used by the randomForest() function in the randomForest package.

**Third step:** Predict the model using our test set

```
predictedRF <- predict(modelRF, test)
```

Now, we predict our model use predict() function, the result of this prediction will contain a prediction value for each data sample in the test data.

**Fourth step:** Create model and Train the model using our training set

```
evaluation = confusionMatrix(predictedRF, test$Churn)
evaluation
```

Now, we evaluate our predictedRF using confusionMatrix() function. This is used to calculate and store the results of the confusion matrix between the prediction generated by the model (predictedRF) and the actual value of the target variable (test\$Churn).

#### Confusion Matrix and Statistics

```

      Reference
Prediction No Yes
No      835  47
Yes     198 986

      Accuracy : 0.8814
      95% CI   : (0.8667, 0.895)
No Information Rate : 0.5
P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.7628

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.8083
      Specificity : 0.9545
      Pos Pred Value : 0.9467
      Neg Pred Value : 0.8328
      Prevalence : 0.5000
      Detection Rate : 0.4042
      Detection Prevalence : 0.4269
      Balanced Accuracy : 0.8814

      'Positive' Class : No
```

From the output above, we can get an information as

- Reference and Prediction:  
Because the 'Positive' Class is No, we can conclude that:
  - **True Negative: 986** indicates the number of correct predictions that the actual data is "No" (not churned or still subscribed) and the model also predicts "No" correctly.
  - **True Positive: 835** indicates the number of correct predictions that the actual data is "Yes" (churned or unsubscribed) and the model also predicts "Yes" correctly.
  - **False Positive: 47** indicates the number of incorrect predictions that the actual data is "No" (not churned or still subscribed), but the model predicts "Yes" (churned or unsubscribed).
  - **False Negative: 198** indicates the number of incorrect predictions that the actual data is "Yes" (churned or unsubscribed), but the model predicts "No" (not churned or still subscribed).
- Performance values:  
**The accuracy value is 0.8814 or 88.14%**, it indicates the proportion of correct predictions out of all predictions.

**The kappa value is 0.7628**, it indicates a good level of agreement between predicted and actual values.

**The sensitivity value is 0.8083**, it indicates that the model has a good ability to correctly identify true positive cases (churned or unsubscribed).

**The specificity value is 0.9545**, it indicates that the model has a good ability to correctly identify true negative cases (non-churned or still subscribed).

**The Pos Pred Value is 0.9467**, it indicates that of all the predictions classified as "No" (non-churned or still subscribed), about 94.67% of them were correct.

**The Neg Pred Value is 0.8328**, it indicates that of all predictions classified as "Yes" (churned or unsubscribed), about 83.28% of them were correct.

**The Prevalence Value is 0.5000**, it indicates that shows that about 50% of the population are positive cases (churned or unsubscribed).

**The Balanced Accuracy is 0.8814**, it is the average of sensitivity and specificity, and shows how well the model can predict overall.

## Conclusion:

Based on the evaluation above, we can conclude that the model was seen to be quite good with high accuracy (**0.8814**) and kappa value (**0.7628**) indicating good agreement between predicted and actual values. Sensitivity (**0.8083**) and specificity (**0.9545**) also showed good ability to correctly identify positive (churned or unsubscribed) and negative (non-churned or still subscribed) cases.

Now, we print the importance of each feature in the models, the output given is the "MeanDecreaseGini" value for each feature in the dataset. The Gini Index is used in the Random Forest algorithm to measure the importance of features in separating target classes. **The greater the mean reduction** in Gini Index, the **more important the feature** is in making predictions. It is important to do this so that we can evaluate, understand, and identify the most influential features in the model.


```
importance(modelRF)
```

	MeanDecreaseGini
gender	84.28351
SeniorCitizen	61.70209
Partner	72.03793
Dependents	66.46954
tenure	592.92141
PhoneService	15.50683
MultipleLines	76.75155
InternetService	156.77820
OnlineSecurity	201.73361
OnlineBackup	94.27275
DeviceProtection	82.57603
TechSupport	167.28826
StreamingTV	70.06521
StreamingMovies	85.11577
Contract	448.69159
PaperlessBilling	80.75710
PaymentMethod	220.05070
MonthlyCharges	533.21698
TotalCharges	595.59368

Based on the output, the tenure feature has a MeanDecreaseGini value of 592.92141, indicating that it is very important in separating different target classes and its importance in making predictions. Similarly, the TotalCharges feature has a MeanDecreaseGini value of 595.59368, indicating its importance in prediction.

Now, we can plot our modelRF using varImpPlot() function, it produces a visual plot showing the relative importance of each feature in our model.

```
varImpPlot(modelRF)
```



Based on the output, we can see the visualization of the importance of each feature which indicates its importance in separating the target class based on MeanDecreaseGini(). We know that the 4 highest MeanDecreaseGini() are TotalCharges, tenure, MonthlyCharges, Contract so we can conclude that these features are very important in making predictions.